

Abhinav Kumar

📍 Bengaluru, India | 📞 +91-7462987886 | ✉️ abkchoubey@gmail.com
🌐 linkedin.com/in/abhinavkumar | 🐙 github.com/abhinavkumar

PROFESSIONAL SUMMARY

Mid-level **Database Developer/Data Engineer** with **3 years** in **healthcare** data infrastructure. Designed **warehouse architectures** and **ETL pipelines** from the ground up, modernized legacy query patterns for performance, and delivered high-impact **ad-hoc analytics** without compromising system reliability. Fluent in **SQL Server and data modeling**; experienced in **identifying root causes of data and performance issues** before they cascade. **Mentor** to junior engineers. Seeking role where I can drive **pipeline architecture, data quality, and scale in a modern stack** (cloud data warehouse, orchestration frameworks, etc.).

WORK EXPERIENCE

Data Engineer/Database Developer

Mirra Healthcare

Bengaluru, India | Sep 2023 – Present

- Designed and maintained **SSIS-based ETL pipelines** to **extract, transform, and load** daily claims data **across 6 client** databases operating on a shared schema, ensuring **consistent and timely data availability** for downstream reporting.
- Automated** a manual claims file import and transformation process, **eliminating 8 hours/week** of manual effort and reducing processing errors.
- Optimized** a critical stored procedure, reducing execution time from **3.5 hours** to under **30 minutes** (86% improvement), significantly improving reporting turnaround.
- Implemented **pre-load data validation controls** (row-count reconciliation, balance checks, duplicate detection) within ETL pipelines, reducing data quality incidents by approximately **40%**.
- Currently build and maintain **5 production ETL pipelines** (3 independently developed) supporting **daily batch data loads** across multiple healthcare clients.
- Designed, developed, and maintained **SSRS reports** and fulfilled **ad-hoc reporting** requests for internal business teams within SLA.
- Reviewed **70+ pull requests** over one year as part of **code quality standards**; deployed changes through **CI/CD pipelines**.
- Mentored 2 junior developers** on T-SQL development, stored procedure design, and SSIS best practices.

TECHNICAL SKILLS

ETL / ELT:	SSIS, Apache Airflow, dbt, CDC-based ingestion, incremental loading, batch pipeline orchestration, medallion architecture
Databases & Warehousing:	SQL Server, PostgreSQL, Databricks SQL / Unity Catalog, star schema, fact/dimension modeling, SCD Type 2, One-Big-Table modeling, BigQuery
Data Quality & Validation:	Pre-load validation, row-count reconciliation, balance checks, duplicate detection, dbt schema tests, singular tests
Reporting & Automation:	SSRS, ad-hoc reporting, manual process automation, TFS, CI/CD pipelines, Git / GitHub, Docker / Docker Compose
Performance Tuning:	Stored procedure optimization, query tuning, ETL pipeline optimization
Collaboration & Leadership:	Code review, pull request review, mentoring junior developers, cross-team reporting support
Cloud Platforms:	AWS (IAM, EC2, S3, Glue, Crawlers, Athena), Azure
Tools:	SQL Server Management Studio (SSMS), Visual Studio Code

PROJECTS

WALMART DATA PIPELINE | Apache Airflow, dbt, Databricks, PostgreSQL

Engineered end-to-end ETL pipeline ingesting retail transactional data from PostgreSQL source via CDC into Databricks, applying medallion architecture transformations (bronze → silver_t → silver_b → gold), and persisting slowly-changing dimensions using dbt snapshots with SCD Type 2.

- Implemented **incremental loading** on six source tables using **dbt materialized='incremental'** with **unique_key** and **timestamp-based change detection** to avoid full refreshes
 - Designed wide **One-Big-Table (OBT) model** in silver_b layer using **Jinja templating** to dynamically generate six-table left-join from configuration dict, **reducing maintenance overhead vs. hand-written SQL**
 - Built **SCD Type 2 dimension tables** (dim_customers, dim_orders, dim_products, etc.) via **dbt snapshots** with **strategy: timestamp** and **dbt_valid_to_current**, capturing row-level history with validity windows
 - Orchestrated **10-task linear DAG** in **Apache Airflow 3.3** using TaskFlow API: **CDC job polling** via **Databricks SDK**, **dbt runs** segmented by layer (silver technical / business / gold), **data quality** tests flagged at **warn severity** to report without blocking
 - Defined **6-table source definitions**, implemented **schema tests** (not_null, unique, accepted_values constraints), and added **singular test** on OBT for null-key detection
- Stack:** Apache Airflow 3.3 | dbt-core 1.12 + dbt-databricks 1.12 | Databricks SQL / Unity Catalog | PostgreSQL | CeleryExecutor | Docker Compose | Python

OLTP-to-Star Data Warehouse | PostgreSQL, dbt, Docker, Git

- Built an **ELT pipeline** that transforms the normalized Pagila DVD-rental **OLTP dataset** into an **analytics-ready star schema** using **dbt and PostgreSQL**.
- **Modeled** reusable **staging** views and **dimensional** tables for customers, films, stores, staff, and dates; created a rental **fact table** with revenue and rental-duration metrics.
- Implemented **SCD Type 2 snapshots** for customer and address records, preserving historical changes for **longitudinal analysis**.
- Added **dbt data-quality tests** for primary-key uniqueness, nullability, and referential integrity; **containerized** the local environment with **Docker Compose**.

EDUCATION

B.Tech in Computer Science & Engineering

Jamshedpur, India | 2018 – 2022

Arka Jain University (JGI) GPA: 8.4 / 10

BEYOND THE KEYBOARD

By day, I **architect clean code**; by night, I chase rhythm across **guitar** strings and **harmonica** reeds. When I'm not at my desk, you'll find me hunting **sixes** on the **cricket** pitch, executing **badminton** drop-shots and **smashes**, or **clutching** late-game wins in **esports**. This balance of intense focus and creative release is what keeps my **problem-solving sharp**.